# Gaurav Chaudhary

gauravmrjatt4@gmail.com | +91-8077598311 | LinkedIn

October 23, 2026

Dear Hiring Manager,

I am applying for the Backend Node.js Developer position listed as remote-first across India. I build production backends in Node.js, ES6+ and TypeScript, and the brief's scope is my daily work: REST and GraphQL API design, schema and query work, JWT and Auth0-backed authentication, webhook processing, queue-backed services, and observability on deployed systems. The posting asks for four to eight years and I have a little over three, all freelance; what I bring against that gap is production scope: a 145-endpoint Express.js and PostgreSQL backend behind a live logistics product, a 107-endpoint tracking platform I have owned since May 2023, and end-to-end ownership of schema, API, queues, tests, deployment and monitoring on both.

Here is how my backend work maps to your responsibilities:

- APIs and schemas: two live backends of 145 and 107 REST endpoints on Express.js and TypeScript, with Prisma schemas of 23 and 12 models on PostgreSQL plus MySQL and MongoDB work, and Apollo GraphQL subscriptions fed by Kafka and Redis pub/sub.
- Queues, caching and distributed systems: RabbitMQ payout queues that keep slow payment jobs off the request path, Redis-cached hot reads, Node.js Cluster workers, and fault-tolerant services for high-concurrency events.
- Money and auth flows: wallet and transaction ledgers, KYC-gated orders, automated payouts and scheduled refunds, Razorpay signature verification on webhooks, and JWT plus Auth0 authorization with role-based access control.
- Testing, CI/CD and operations: an automated unit and integration suite in production, Docker, GitHub Actions pipelines, and Prometheus, Grafana, Loki and Sentry monitoring of live incidents.

On tools the posting names that I lack: my AWS work is at EC2 level, so Lambda, API Gateway, S3, CloudWatch and DynamoDB are not in my production history, and my queue experience is RabbitMQ and Kafka rather than SQS. I have not used ELK, Datadog, GitLab CI, or Jira and Confluence; my equivalents are Loki, Prometheus, Grafana and Sentry, GitHub Actions, and direct client communication instead of an Agile board. The posting names Jest and supertest; mine is a ten-file suite in a live backend. I have not mentored inside a team, but I have kept a live codebase test-covered for over three years for clients who read it before release.

Remote delivery for distributed clients since May 2023 has made async updates, release notes, and owning incidents after deployment routine. I would welcome the chance to walk through the schema, queue and query decisions behind either backend.

Kind regards,

Gaurav Chaudhary